

# CSCI 403: Databases

## 4 - Basic DDL

### Schema

In the relational model of the database, we used schema as a term to describe the names and attributes of relations and the constraints on the relations in a database. In the SQL standard, schema is defined more broadly as a container for database objects, including tables, constraints, views, indexes, etc. A SQL database can contain many schemas.

- Like namespaces, a schema allows us to separate logical units in a database and avoid name collisions.
- Schemas also makes for easier application of security policies by letting us grant permissions on whole schemas at a time rather than table by table.
- Schemas are generally something created and managed by the database administrator/owner, not by the database user or programmer. (Acknowledging that of course, programmers may also be DBAs or DB owners!)
- You do not (AFAIK) have rights to create schemas in the course database. You have your own schema, named by your Mines id. Anything you create by default will be created in your personal schema. You also have read-only access to a schema named "public".

### Catalog

In SQL, the *catalog* can be viewed as either a higher-level container containing schemas, or as the set of system tables within the database describing all of the other objects in the database. All relational DBMSes have some form of system

tables which can be read to determine the structure of objects in the database. To see the system tables in `psql`, enter the command "`\dS`" at the prompt.

### Table Creation

SQL uses a very "English-like" syntax. To create a table in SQL, you issue a `CREATE TABLE` command. A somewhat abbreviated syntax of the command looks like:

```
CREATE TABLE [schemaname.]tablename (  
  { columnname datatype [NOT NULL] [UNIQUE]  
  [PRIMARY KEY] | table_constraint } [...])
```

Which probably looks like gibberish... you will need to learn how to read SQL documentation, most of which looks like this. Briefly, anything in `[]` is optional, while the construct `{ A | B }` means you can have A or B. If you see `[,...]`, that means you can have more copies of the last optional thing (with a comma separator). Anything else is required.

Here's an example of creating a throwaway table, showing some of the different data types you might use:

```
CREATE TABLE yourid.stuff (  
  id serial PRIMARY KEY,  
  name text NOT NULL,  
  age integer,  
  gender char(1),  
  salary numeric(9,2),  
  favorite_constant double precision,  
  date_hired date);
```

This will create a table named "stuff" in the schema "yourid" with several columns, each of a different type. (Note that text and serial types are specific to PostgreSQL, although similar types exist in other DBMSes.) This construction sets the column "id" as a primary key for the table (just the one column in the key), and the name

field is constrained to not contain NULL values. We can also create foreign key references or other constraints within the table creation statement. Table constraints can also be created in a separate declaration in the comma-separated list within the (); here's the same table, but with a two-column primary key (on name and age):

```
CREATE TABLE yourid.stuff (
  id serial,
  name text NOT NULL,
  age integer,
  gender char(1),
  salary numeric(9,2),
  favorite_constant double precision,
  date_hired date,
  PRIMARY KEY (name, age));
```

## Types

SQL defines a number of types for attributes in your relations, and most DBMSes define additional types. Some types that you will find useful are listed below.

### *Integers*

- INTEGER - 32-bit integers
- SMALLINT - 16-bit integers
- BIGINT - 64-bit integers

### *Fixed-precision numeric (exact)*

- NUMERIC(w,p) - Defines numbers with a maximum of w digits, and a precision of 2
- DECIMAL(w,p) - same as NUMERIC(w,p)

### *Floating point (inexact)*

- REAL - 32-bit floating point
- DOUBLE PRECISION - 64-bit floating point

### *Strings*

- CHAR(n) - strings of length exactly n, padded with spaces if necessary
- VARCHAR(n) - variable length strings of max length n
- TEXT - variable length strings, no limit (PostgreSQL type)

### *Date/Time*

- DATE - holds dates. You can enter DATE values as strings in the format 'YYYY-MM-DD', other formats maybe possible depending on DBMS.
- TIME - holds times. Format 'HH:MM:SS', can also add decimal points after for sub-second times. Optionally, timezone can also be included.
- TIMESTAMP - date and time.

### *Other types*

- BOOLEAN - holds true/false values - various formats are compatible
- SERIAL - an auto-incrementing integer type (PostgreSQL type)
- MONEY
- ... many more

## CREATE... AS

An easy way to create a table from a SELECT query; do

```
CREATE TABLE schemaname.tablename AS
SELECT ... ;
```

This will create a table with attributes and types determined by the SELECT query result. Column renaming together with functions, joins, etc., makes this a powerful way to create a new table. Note, however, that this does not create any keys or constraints - these will have to be added on using ALTER TABLE.

## Notes on Workflow

Tables/schemas are generally created only occasionally and modified (relatively) seldom thereafter. [You may have a different experience in an agile environment; one reason NoSQL is popular right now is the "schema-less" nature of NoSQL databases.]

When making tables, it is easy to make small mistakes or wish you had done something differently. Changes after the fact are often harder than simply dropping everything and starting over. Thus: make scripts! Your scripts should (optionally) drop everything, then create all of your tables and constraints, load all of your data,

etc. Re-run until you are happy with the result.  
(Having these scripts will help a lot downstream  
when your application goes into development for  
version 2...)