

Query Optimization

1. Conjunction of SELECT operations \rightarrow cascade of selects
$$\sigma_{c1} AND \sigma_{c2} AND \dots$$
$$= \sigma_{c1}(\sigma_{c2}(\sigma_{c3}(\dots)))$$
2. Commutativity of SELECT
$$\sigma_{c1}(\sigma_{c2}(R)) = \sigma_{c2}(\sigma_{c1}(R))$$
3. Cascade of Projections operations: only last projection step counts
$$\pi_x(\pi_y(\pi_z(\dots))) = \pi_x(\dots)$$
4. Commutativity of σ, π
5. \bowtie and X are commutative
6. Commutativity of σ with \bowtie and X when σ condition applies to a single table involved in the \bowtie or X
$$\sigma_{A[x]=\dots} (A \times B) = (\sigma_{A[x]=\dots}(A) \times B)$$
7. Commutativity of π with \bowtie and X when join (or X) only involved attributes in projection
8. Associativity of \bowtie, X

Heuristic Algorithm for Optimizing Query Tree

1. Break up any conjunction of select conditions into a cascade
- 2.

Transactions

- Problem: multiple users accessing/modifying db at the same time
- Basic model:
 - Read an object
 - Write an object (e.g. blocks, rows, tables)

Concurrency Problems

1. Lost updates
2. Dirty read

Improve Model:

- Begin transaction
- Read/write
- Commit
- Rollback
- Goal: transaction treated as atomic unit of work ("all or nothing")
- Successful commit \rightarrow